

Daizhi Liao

Undergraduate researcher · Reinforcement Learning & Embodied AI · PhD applicant Fall 2027

Guangzhou, China · ldz@e.gzhu.edu.cn · dafahaha.github.io · github.com/dafahaha

EDUCATION

Guangzhou University — B.Eng., Network Engineering

Expected 2027

Third-year undergraduate. Focus: reinforcement learning, embodied AI, and the empirical auditing of deployed AI systems.

RESEARCH & PROJECTS

TransitTruth — Behavioral-Fingerprint Audit of AI API Relays

2026

Solo design & implementation. Open-source tool that verifies whether an AI API relay actually serves the advertised model, by probing statistical behavior (random-number, token, latency distributions) against baselines with chi-square / KS tests and Bayesian updating. Built a zero-install, fully client-side web demo and a 20-page technical report validating same-family discrimination: gpt-4o vs gpt-4o-mini, TVD up to 0.90, 91.3% in-distribution resubstitution posterior (not held-out).

Code & demo: github.com/dafahaha/transit-truth · live demo

embodied-offpolicy-study — Pre-registered Three-Arm SAC Study on Hopper-v4

2026

Soft Actor-Critic implemented from scratch in PyTorch (no Stable-Baselines3), run as a pre-registered three-arm controlled study on Hopper-v4 (CPU, 120k env steps, 3 seeds/arm): Baseline (reward scale 1.0) vs RewardScaled (0.1) vs FixedAlpha (0.1, automatic temperature off, α fixed at 0.2). Asymptotic return 1085 ± 790 / 1118 ± 705 / 1404 ± 715 ; pairwise Welch $p \in [0.55, 0.95]$, one-way ANOVA $F \approx 0.22$, $p \approx 0.90$ — no statistically distinguishable effect detected (low power at $n=3$). At 60k the reward-scaled arm looked $\sim 3\times$ better, but extending to the pre-registered 120k budget and adding the fixed-temperature control overturned that apparent gain as early seed noise: an honest pre-registered null. Raw CSVs committed, curves reproduced by script.

Code: github.com/dafahaha/embodied-offpolicy-study

rl-deploy-bench — Cross-Platform RL Deployment & Benchmarking

v1.1.0

Config-driven CLI: export Stable-Baselines3 policies to ONNX/TorchScript, build TensorRT engines with FP16/INT8 quantization, and benchmark latency, throughput, and accuracy across x86 GPU, NVIDIA Jetson, and CPU, with auto-generated HTML reports. A reproducible harness for the accuracy-latency tradeoff of quantized RL policies.

OPEN SOURCE — MERGED CONTRIBUTIONS

- **microsoft/TextWorld #377** (merged) — Fixed `env.step` crash on multi-command input; non-greedy score parsing plus regression test.
- **shmuma/ptan #57** (merged) — Fixed Gymnasium API compatibility in experience sources.
- **Algorineko/AgenticArXiv-RL #72** (merged) — Added MIT License, CONTRIBUTING guide, issue templates, and CI.
- Additional review-phase bug fixes upstream in vLLM, pandas, google/brax, Farama-Foundation/PettingZoo, and Hugging Face datasets & diffusers.

TECHNICAL SKILLS

RL & learning: PyTorch, Stable-Baselines3, Gymnasium, MuJoCo | **Deployment:** ONNX, ONNX Runtime, TensorRT, TorchScript, INT8/FP16 quantization | **Systems:** Python, C++, Linux, Docker, Git, CI/CD, CUDA